

# Attention via $\log \sum \exp$ Energy

Alexander Tschantz

January 21, 2025

## 1 General Framework

**Setup.** We consider a single set of nodes  $\mathbf{v} = \{\mathbf{v}_a : a \in \{1, 2, \dots, A\}\}$ , where each node  $\mathbf{v}_a \in \mathbb{R}^d$ . The relationships between these nodes are defined by a set of  $M$  energy functions  $\{E_m : m \in \{1, 2, \dots, M\}\}$ . Each energy function  $E_m$  defines a subset of nodes acting as *children*  $C_m \subseteq \{1, 2, \dots, A\}$  and a subset acting as *parents*  $P_m \subseteq \{1, 2, \dots, A\}$ , which may overlap.

**Energy.** Each energy function  $E_m$  defines a similarity function:

$$\text{sim}(\mathbf{v}_c, \mathbf{v}_p) \quad : \quad \mathbb{R}^d \times \mathbb{R}^d \rightarrow \mathbb{R}, \quad (1)$$

which produces a scalar similarity between a child  $\mathbf{v}_c$  and a parent  $\mathbf{v}_p$ . Using  $\{\mathbf{v}_c\} = \{\mathbf{v}_c : c \in C_m\}$  and  $\{\mathbf{v}_p\} = \{\mathbf{v}_p : p \in P_m\}$ , the energy for  $E_m$  is defined as:

$$E_m(\{\mathbf{v}_c\}, \{\mathbf{v}_p\}) = - \sum_{c \in C_m} \ln \left( \sum_{p \in P_m} \exp(\text{sim}(\mathbf{v}_c, \mathbf{v}_p)) \right). \quad (2)$$

The global energy sums over all energy functions:

$$E(\{\mathbf{v}\}) = \sum_{m=1}^M E_m(\{\mathbf{v}_c\}, \{\mathbf{v}_p\}). \quad (3)$$

**Gradient Updates.** For a single node  $\mathbf{v}_a$ , the gradient of the global energy  $E$  w.r.t.  $\mathbf{v}_a$  decomposes into two terms. Let  $\mathcal{M}_c(a) = \{m : a \in C_m\}$  denote the energy functions where  $\mathbf{v}_a$  acts as a *child*, and  $\mathcal{M}_p(a) = \{m : a \in P_m\}$  the energy functions where  $\mathbf{v}_a$  acts as a *parent*. Then:

$$\begin{aligned} -\frac{\partial E}{\partial \mathbf{v}_a} = & \underbrace{\sum_{m \in \mathcal{M}_c(a)} \sum_{p \in P_m} \text{softmax}_p(\text{sim}(\mathbf{v}_a, \mathbf{v}_p)) \frac{\partial}{\partial \mathbf{v}_a} \text{sim}(\mathbf{v}_a, \mathbf{v}_p)}_{\mathbf{v}_a \text{ acting as a child}} \\ & + \underbrace{\sum_{m \in \mathcal{M}_p(a)} \sum_{c \in C_m} \text{softmax}_a(\text{sim}(\mathbf{v}_c, \mathbf{v}_a)) \frac{\partial}{\partial \mathbf{v}_a} \text{sim}(\mathbf{v}_c, \mathbf{v}_a)}_{\mathbf{v}_a \text{ acting as a parent}}. \end{aligned} \quad (4)$$

The first term captures contributions from  $\mathbf{v}_a$  being explained by its parents, while the second term captures contributions from  $\mathbf{v}_a$  explaining its children.

## 2 Gaussian Mixture Models (GMMs)

**Setup.** We have  $N$  data points (children)  $\mathbf{x}_i \in \mathbb{R}^d$ ,  $i \in C = \{1, \dots, N\}$ , and  $K$  mixture components (parents), each with mean  $\boldsymbol{\mu}_k \in \mathbb{R}^d$  and covariance  $\boldsymbol{\Sigma}_k$ ,  $k \in P = \{1, \dots, K\}$ . Let  $\pi_k$  be the mixing proportion.

**Similarity function.** We define

$$\text{sim}(\mathbf{x}_i, \boldsymbol{\mu}_k) = \ln \pi_k - \frac{1}{2} (\mathbf{x}_i - \boldsymbol{\mu}_k)^\top \boldsymbol{\Sigma}_k^{-1} (\mathbf{x}_i - \boldsymbol{\mu}_k).$$

**Energy.**

$$E^{\text{GMM}}(\{\mathbf{x}_i\}, \{\boldsymbol{\mu}_k\}) = - \sum_{i=1}^N \ln \left( \sum_{k=1}^K \exp(\text{sim}(\mathbf{x}_i, \boldsymbol{\mu}_k)) \right). \quad (5)$$

**Gradients.** If we differentiate w.r.t.  $\boldsymbol{\mu}_k$ , then

$$-\frac{\partial E^{\text{GMM}}}{\partial \boldsymbol{\mu}_k} = \sum_{i=1}^N \text{softmax}_k(\mathbf{A}_{ik}) \boldsymbol{\Sigma}_k^{-1}(\mathbf{x}_i - \boldsymbol{\mu}_k).$$

Setting this gradient to zero yields the usual GMM M-step:

$$\boldsymbol{\mu}_k = \frac{\sum_{i=1}^N \text{softmax}_k(\mathbf{A}_{ik}) \mathbf{x}_i}{\sum_{i=1}^N \text{softmax}_k(\mathbf{A}_{ik})}.$$

### 3 Cross Attention

**Setup.** We have a set of child vectors (queries)  $\mathbf{Q} \in \mathbb{R}^{d \times N_Q}$  and a set of parent vectors (keys)  $\mathbf{K} \in \mathbb{R}^{d \times N_K}$ . Let

$$C = \{1, \dots, N_Q\}, \quad P = \{1, \dots, N_K\},$$

so  $\mathbf{v}_c = \mathbf{q}_c$  is the  $c$ -th query, and  $\mathbf{v}_p = \mathbf{k}_p$  is the  $p$ -th key. Suppose we have learnable weight matrices  $\mathbf{W}^Q, \mathbf{W}^K \in \mathbb{R}^{d \times d}$ . Then

$$\mathbf{q}_c = \mathbf{W}^Q \mathbf{x}_c^Q, \quad \mathbf{k}_p = \mathbf{W}^K \mathbf{x}_p^K,$$

where  $\mathbf{x}_c^Q$  is the raw  $c$ -th query token and  $\mathbf{x}_p^K$  the raw  $p$ -th key token.

**Similarity function.**

$$\text{sim}(\mathbf{q}_c, \mathbf{k}_p) = \mathbf{q}_c^\top \mathbf{k}_p.$$

**Energy.**

$$E^{\text{Cross}}(\{\mathbf{q}_c\}, \{\mathbf{k}_p\}) = - \sum_{c=1}^{N_Q} \ln \left( \sum_{p=1}^{N_K} \exp(\mathbf{q}_c^\top \mathbf{k}_p) \right). \quad (6)$$

**Gradients.**

$$-\frac{\partial E^{\text{Cross}}}{\partial \mathbf{q}_c} = \sum_{p=1}^{N_K} \text{softmax}_p(\mathbf{q}_c^\top \mathbf{k}_p) \mathbf{k}_p. \quad (7)$$

$$-\frac{\partial E^{\text{Cross}}}{\partial \mathbf{k}_p} = \sum_{c=1}^{N_Q} \text{softmax}_p(\mathbf{q}_c^\top \mathbf{k}_p) \mathbf{q}_c. \quad (8)$$

When mapping back to the raw tokens  $\mathbf{x}_c^Q$  or  $\mathbf{x}_p^K$ , chain-rule multiplies by  $\mathbf{W}^Q$  or  $\mathbf{W}^K$ , respectively.

### 4 Hopfield Networks

**Setup.** We have a set of *children* data vectors  $\mathbf{x}_i \in \mathbb{R}^d, i \in C = \{1, \dots, N\}$ , and a set of *parent* memory vectors  $\mathbf{m}_\mu \in \mathbb{R}^d, \mu \in P = \{1, \dots, K\}$ .

**Similarity function.**

$$\text{sim}(\mathbf{x}_i, \mathbf{m}_\mu) = \mathbf{x}_i^\top \mathbf{m}_\mu.$$

**Energy.**

$$E^{\text{Hopfield}}(\{\mathbf{x}_i\}, \{\mathbf{m}_\mu\}) = - \sum_{i=1}^N \ln \left( \sum_{\mu=1}^K \exp(\mathbf{x}_i^\top \mathbf{m}_\mu) \right). \quad (9)$$

**Gradients.**

$$-\frac{\partial E^{\text{Hopfield}}}{\partial \mathbf{x}_i} = \sum_{\mu=1}^K \text{softmax}_{\mu}(\mathbf{x}_i^{\top} \mathbf{m}_{\mu}) \mathbf{m}_{\mu}. \quad (10)$$

$$-\frac{\partial E^{\text{Hopfield}}}{\partial \mathbf{m}_{\mu}} = \sum_{i=1}^N \text{softmax}_{\mu}(\mathbf{x}_i^{\top} \mathbf{m}_{\mu}) \mathbf{x}_i. \quad (11)$$

## 5 Slot Attention

**Setup.** Let  $\mathbf{x}_j \in \mathbb{R}^d$ ,  $j \in C = \{1, \dots, N\}$  be the children (tokens), and  $\boldsymbol{\mu}_i \in \mathbb{R}^d$ ,  $i \in P = \{1, \dots, S\}$  be the parents (slots). We typically apply linear transforms  $\mathbf{W}_K, \mathbf{W}_Q \in \mathbb{R}^{d \times d}$  to form

$$\text{sim}(\mathbf{x}_j, \boldsymbol{\mu}_i) = (\mathbf{W}_K \mathbf{x}_j)^{\top} (\mathbf{W}_Q \boldsymbol{\mu}_i).$$

**Energy.**

$$E^{\text{Slot}}(\{\mathbf{x}_j\}, \{\boldsymbol{\mu}_i\}) = -\sum_{j=1}^N \ln \left( \sum_{i=1}^S \exp(\text{sim}(\mathbf{x}_j, \boldsymbol{\mu}_i)) \right). \quad (12)$$

**Gradients.**

$$-\frac{\partial E^{\text{Slot}}}{\partial \mathbf{x}_j} = \sum_{i=1}^S \text{softmax}_i(\text{sim}(\mathbf{x}_j, \boldsymbol{\mu}_i)) \mathbf{W}_K^{\top} \mathbf{W}_Q \boldsymbol{\mu}_i. \quad (13)$$

$$-\frac{\partial E^{\text{Slot}}}{\partial \boldsymbol{\mu}_i} = \sum_{j=1}^N \text{softmax}_i(\text{sim}(\mathbf{x}_j, \boldsymbol{\mu}_i)) \mathbf{W}_Q^{\top} \mathbf{W}_K \mathbf{x}_j. \quad (14)$$

## 6 Self-Attention

**Setup.** In self-attention, every node can act as both a child (query) and a parent (key). Concretely, let us have  $N$  tokens  $\{\mathbf{x}_1, \dots, \mathbf{x}_N\}$ . We form

$$\mathbf{q}_i = \mathbf{W}^Q \mathbf{x}_i, \quad \mathbf{k}_i = \mathbf{W}^K \mathbf{x}_i,$$

for  $i = 1, \dots, N$ . Thus the set  $C = \{1, \dots, N\}$  and  $P = \{1, \dots, N\}$  coincide, with

$$\text{sim}(\mathbf{x}_c, \mathbf{x}_p) = (\mathbf{W}^Q \mathbf{x}_c)^{\top} (\mathbf{W}^K \mathbf{x}_p).$$

**Energy.**

$$E^{\text{SA}}(\{\mathbf{x}_i\}) = -\sum_{c=1}^N \ln \left( \sum_{p=1}^N \exp \left( (\mathbf{W}^Q \mathbf{x}_c)^{\top} (\mathbf{W}^K \mathbf{x}_p) \right) \right). \quad (15)$$

**Gradients.** Since each  $\mathbf{x}_i$  is *both* a child and a parent, its gradient is a sum of two terms (the child side and the parent side). Writing it out explicitly:

$$\begin{aligned} -\frac{\partial E^{\text{SA}}}{\partial \mathbf{x}_i} &= \underbrace{\sum_{p=1}^N \text{softmax}_p \left( (\mathbf{W}^Q \mathbf{x}_i)^{\top} (\mathbf{W}^K \mathbf{x}_p) \right) \mathbf{W}_Q^{\top} \mathbf{W}_K \mathbf{x}_p}_{\text{child } i \text{ being explained by parents } p} \\ &\quad + \underbrace{\sum_{c=1}^N \text{softmax}_i \left( (\mathbf{W}^Q \mathbf{x}_c)^{\top} (\mathbf{W}^K \mathbf{x}_i) \right) \mathbf{W}_K^{\top} \mathbf{W}_Q \mathbf{x}_c}_{\text{parent } i \text{ explaining children } c}. \end{aligned} \quad (16)$$